

CSCB09 2025 Summer Assignment 3

Due: July 27 2025, 11:59PM

In this assignment, you will practice making child processes, exec'ing other programs, and setting up I/O redirections and/or pipes.

As usual, you should aim for reasonably efficient algorithms and reasonably organized, comprehensible code.

Process Substitution

bash has a convenient extra feature “process substitution” not in sh. In the simplest form it goes like

```
command1 <(command2)
```

command1 will see a command line argument. We users don't see it, but command1 does. Let's call it F because I need to refer to it later. command1 can treat F as a filename and open it for reading (only); command2 can write data to stdout. Here is the magical part: Whenever command1 reads from “file” F, it sees command2's output!

But it can be more general!

```
command1 <(command2a) <(command2b)
```

The above means: There are secret filenames F2a and F2b. When command1 reads from F2a, that's command2a's output; when command1 reads from F2b, that's command2b's output.

```
command1 <(command2 <(command3))
```

The above means: There are secret filenames F and G. When command1 reads from F, that's command2's output; when command2 reads from G, that's command3's output.

Many students used this feature to evade learning a required topic. The prof is not pleased. The prof's revenge is to make you implement this feature so you have a harder topic to learn!

(On the bright side, haven't you always wondered how bash pulls off this trick?)

Discovery (optional, not graded)

Write a shell script or C program that prints the command line arguments it sees. It doesn't actually need to read any file. Use it as command1. This is to discover the secret F to help find out what's going on. Example bash command:

```
./yourprog <(echo hello)
```

What do you find? What do you think it means?

Theory

/dev/fd is a special directory (not on real disk) made up by the kernel. Whenever a process P looks into it, P finds numeric filenames, and the numbers are exactly P's current open file descriptors. And if P reads from, say, /dev/fd/42, it's equivalent to reading from file descriptor 42.

(Naturally, different processes see different content, since the kernel is making it up on the fly, tailor-made to the process that asks.)

With this friendly service from the kernel, bash can pull off its trick. (It is still a pipe.) We assume that commands are executable programs with arguments in this assignment, even though bash is much more flexible than that.

For example, here is how to do `command1 <(command2)`:

- Create a pipe. Let's say the read-end FD is 42, write-end FD is 43, for the sake of concreteness, but this is just an example.
- Create a child process for command1 with argument `/dev/fd/42` (the secret F above). So if command1 opens that "file" and read, it effectively reads from the pipe. Note that this is **not** stdin redirection.
- Create a child process for command2 with stdout redirected as FD 43. So if command2 writes to stdout, the data goes to the pipe.
- All 3 processes should also close certain unneeded file descriptors, as discussed in the lecture concerning pipe hygiene. It is your job to figure out which and when.

Implementation

The function you will implement (create and hand in probsub.c) is:

```
int run(struct command *cmd, int *wstatus);
```

`struct command` specifies a command; see probsub.h, but informally it has: program name, number of arguments, and (pointer to) the arguments.

`struct argument` specifies an argument, see probsub.h; informally it has two cases: a string, or a command (meaning that process substitution is to be done). We will only implement the `<(cmd)` kind of substitutions.

Further clarifications can be found as examples in sample-main.c with comments stating the bash equivalents. In particular, if the bash equivalent is of the form `progname foo`, then the `numargs` field is 1, and the `args` array has just one element (for `foo`); so this is different from what `exec*`() expects, and you have to bridge the gap.

`run()` should run the given command in a child process, wait for it to terminate, and store its wait status at the address given by `wstatus` (you may assume that it is non-NUL). Needless to say, extra children are also necessary for arguments of the `SUBST` case; for simplicity, you do not need to call `wait()` or `waitpid()` on those other children.

`run()` should return 0 except under the conditions in the next section.

Error handling

Any process: If `exec*`() fails, exit with exit code 127. Error message is up to you, but send any to stderr.

Other errors: You may ignore or handle, your choice. If you choose to handle: 1. send any error message to stderr; 2. if the error occurs in the parent, `run()` returns -1; 3. if the error occurs in a child, exit with exit code 127.